

predicting-health-insurance-price

May 10, 2025

1 Predicting Health Insurance Price for an individual or family

2 Target Variable (Dependent Variable):

- health_insurance_price

3 Data Understanding

```
[360]: # import necessary libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

In this section, we import essential Python libraries for:

- Data handling (Pandas, NumPy)
- Visualization (Seaborn, Matplotlib)
- Preprocessing and building machine learning models (scikit-learn)

```
[361]: # Load the dataset
df = pd.read_excel("/content/Health_insurance_cost_dataset.xlsx")
print(df)
```

	age	gender	BMI	Children	smoking_status	location	\
0	19.0	female	NaN	0	yes	southwest	
1	18.0	male	33.770	1	no	southeast	
2	28.0	male	33.000	3	no	southeast	
3	33.0	male	22.705	0	no	northwest	
4	32.0	male	28.880	0	no	northwest	
...	
1333	50.0	male	30.970	3	no	northwest	
1334	18.0	female	31.920	0	no	northeast	
1335	18.0	female	36.850	0	no	southeast	
1336	21.0	female	25.800	0	no	southwest	
1337	61.0	female	29.070	0	yes	northwest	

health_insurance_price

```

0          16884.92400
1          1725.55230
2          4449.46200
3          21984.47061
4          3866.85520
...
1333       10600.54830
1334       2205.98080
1335       1629.83350
1336       2007.94500
1337       29141.36030

```

[1338 rows x 7 columns]

We load the dataset using Pandas, then explore its structure and check for:

- Null values
- Duplicate records
- Data types
- Summary statistics

```

[362]: # Display first 5 rows
print("First 5 rows of the dataset: ")
print(df.head())

```

First 5 rows of the dataset:

	age	gender	BMI	Children	smoking_status	location \
0	19.0	female	NaN	0	yes	southwest
1	18.0	male	33.770	1	no	southeast
2	28.0	male	33.000	3	no	southeast
3	33.0	male	22.705	0	no	northwest
4	32.0	male	28.880	0	no	northwest

	health_insurance_price
0	16884.92400
1	1725.55230
2	4449.46200
3	21984.47061
4	3866.85520

Its show the first 5 rcords from the dataset

```

[363]: # Display last 5 rows
print("Last 5 rows of the dataset: ")
print(df.tail())

```

Last 5 rows of the dataset:

	age	gender	BMI	Children	smoking_status	location \
1333	50.0	male	30.97	3	no	northwest
1334	18.0	female	31.92	0	no	northeast

```

1335  18.0  female  36.85          0          no  southeast
1336  21.0  female  25.80          0          no  southwest
1337  61.0  female  29.07          0          yes  northwest

```

```

health_insurance_price
1333          10600.5483
1334          2205.9808
1335          1629.8335
1336          2007.9450
1337          29141.3603

```

Its show the last 5 rcords from the dataset

```

[364]: # size of the dataset
print("Dataset Shape:", end=" ")
print(df.shape) # (rows, columns)

```

Dataset Shape: (1338, 7)

Its show shape of the dataset like how many rows and columns are present in the datssset.

```

[365]: # Check data types and memory usage
print("Dataset Information:")
print(df.info())

```

```

Dataset Information:
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1338 entries, 0 to 1337
Data columns (total 7 columns):
#   Column                Non-Null Count  Dtype
---  -
0   age                   1310 non-null   float64
1   gender                1338 non-null   object
2   BMI                   1315 non-null   float64
3   Children              1338 non-null   int64
4   smoking_status        1338 non-null   object
5   location              1338 non-null   object
6   health_insurance_price 1336 non-null   float64
dtypes: float64(3), int64(1), object(3)
memory usage: 73.3+ KB
None

```

The dataset has 1,338 entries and 7 columns, with missing values in age, BMI, and health_insurance_price, while all other columns are complete.

And are three catgorical columns - gender, smoking_status, location.

```

[366]: # Check for missing values
print("Missing values in each column:")
print(df.isnull().sum())

```

Missing values in each column:

```
age                28
gender             0
BMI                23
Children           0
smoking_status     0
location           0
health_insurance_price  2
dtype: int64
```

There are missing values in age, BMI, and health_insurance_price.

```
[367]: # Summary statistics for numerical columns
print("Summary Statistics:")
print(df.describe())
```

Summary Statistics:

	age	BMI	Children	health_insurance_price
count	1310.000000	1315.000000	1338.000000	1336.000000
mean	39.166412	30.638217	1.094918	13268.527719
std	14.055378	6.110302	1.205493	12112.797724
min	18.000000	15.960000	0.000000	1121.873900
25%	26.000000	26.210000	0.000000	4744.325050
50%	39.000000	30.305000	1.000000	9382.033000
75%	51.000000	34.580000	2.000000	16604.302645
max	64.000000	53.130000	5.000000	63770.428010

```
[368]: # Summary statistics for categorical columns
print("Categorical feature summary:")
print(df.describe(include=['object']))
```

Categorical feature summary:

	gender	smoking_status	location
count	1338	1338	1338
unique	2	2	4
top	male	no	southeast
freq	676	1064	364

```
[369]: # Unique values in each column
print("Unique values per column:")
for col in df.columns:
    print(f"{col}: {df[col].nunique()} unique values")
```

Unique values per column:

```
age: 47 unique values
gender: 2 unique values
BMI: 545 unique values
Children: 6 unique values
smoking_status: 2 unique values
```

location: 4 unique values
health_insurance_price: 1335 unique values

```
[370]: # Check for duplicate rows
print("Number of Duplicate Rows:",end=" ")
print(df.duplicated().sum())
```

Number of Duplicate Rows: 1

```
[371]: # Show the duplicate rows
print("Duplicate Rows:")
print(df[df.duplicated()])
```

Duplicate Rows:

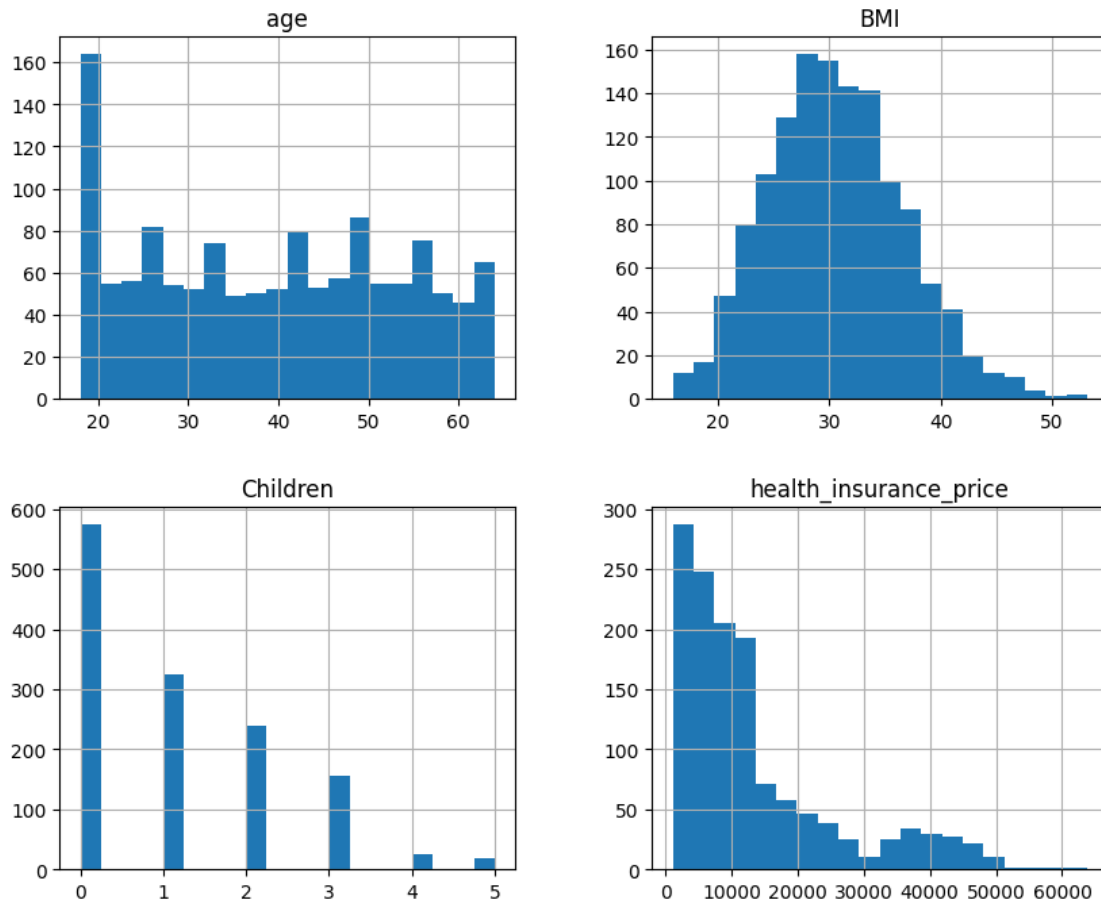
	age	gender	BMI	Children	smoking_status	location	\
581	19.0	male	30.59	0	no	northwest	

	health_insurance_price
581	1639.5631

Its show duplicated row in the dataset.

```
[372]: # Histogram for numerical columns
df.hist(figsize=(10, 8), bins=20)
plt.suptitle("Feature Distributions", fontsize=16)
plt.show()
```

Feature Distributions

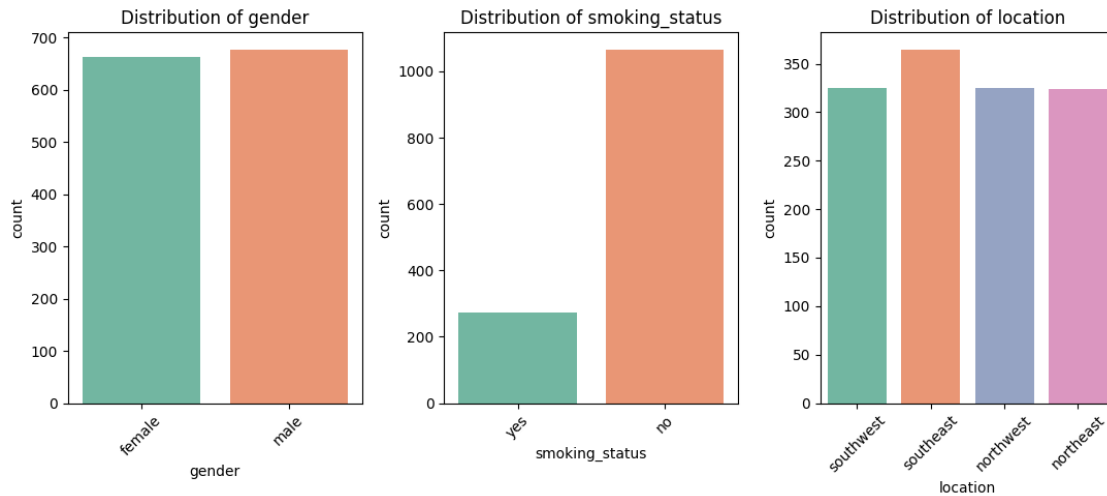


Its show distribution of all numerical columns

```
[373]: # Countplot for categorical variables
categorical_features = ["gender", "smoking_status", "location"]
plt.figure(figsize=(11,5))

for i, col in enumerate(categorical_features, 1):
    plt.subplot(1, len(categorical_features), i)
    sns.countplot(data=df, x=col, hue=col, palette="Set2", legend=False) # Use
    ↪ 'first_sheet_df'
    plt.title(f"Distribution of {col}")
    plt.xticks(rotation=45)

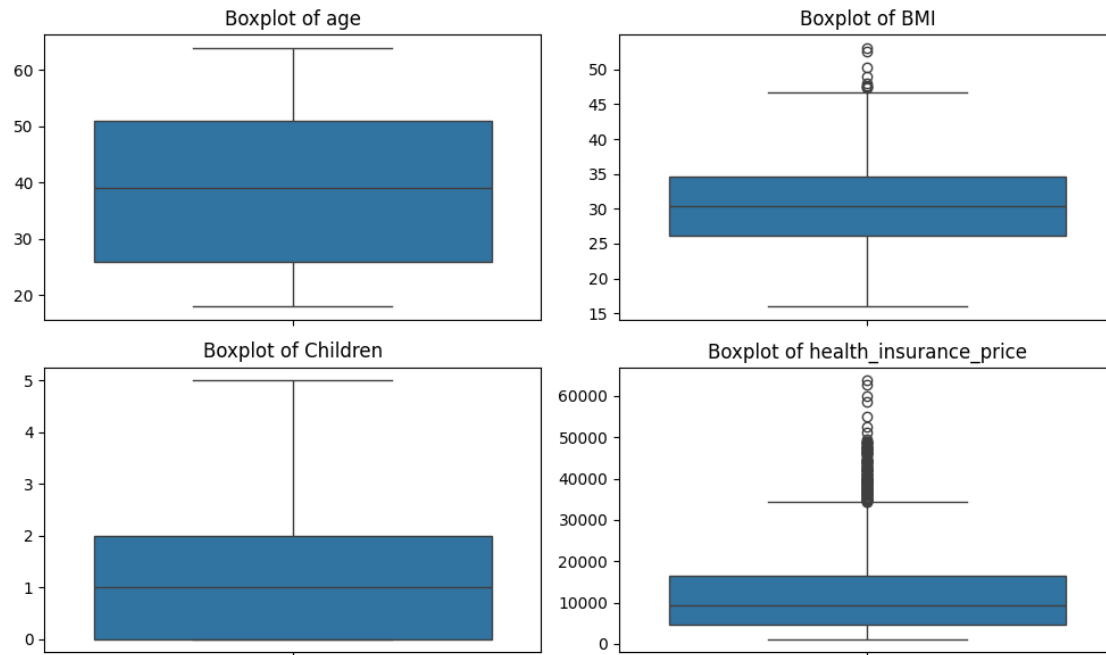
plt.tight_layout()
plt.show()
```



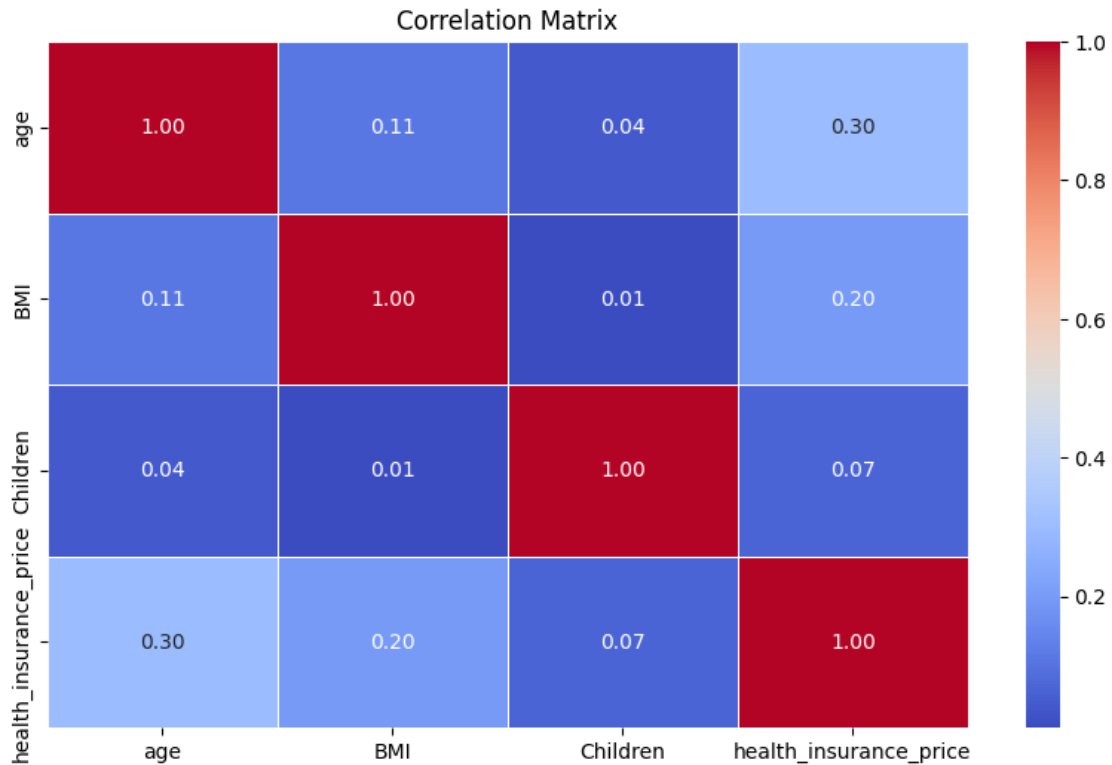
Its show distribution of all categorical columns

```
[374]: # Boxplot for numerical features to check for outliers
numerical_features = ['age', 'BMI', 'Children', 'health_insurance_price']

plt.figure(figsize=(10,6))
for i, col in enumerate(numerical_features, 1):
    plt.subplot(2,2,i)
    sns.boxplot(data=df, y=col)
    plt.title(f"Boxplot of {col}")
    plt.ylabel("")
plt.tight_layout()
plt.show()
```



```
[375]: # # Plot heatmap to visualize correlation between numeric features
plt.figure(figsize=(10,6))
sns.heatmap(df.select_dtypes(include=['number']).corr(), annot=True,
            cmap="coolwarm",fmt=".2f",linewidths=0.5)
plt.title("Correlation Matrix")
plt.show()
```

We visualize the correlation between features using a heatmap to understand relationships and multicollinearity.

4 Data Preprocessing

4.1 Handling Duplicates

```
[376]: # check for duplicate rows
duplicate_count = df.duplicated().sum()
print(f"Number of duplicate rows: {duplicate_count}")
```

Number of duplicate rows: 1

```
[377]: # Remove duplicate rows
df = df.drop_duplicates()
```

```
[378]: # Verify removal
print(f"New dataset shape after removing duplicates: {df.shape}")
```

New dataset shape after removing duplicates: (1337, 7)

5 Handling Missing Values

```
[379]: # Check missing values before imputation
print("Missing values before imputation: ")
print(df.isnull().sum())
```

Missing values before imputation:

```
age          28
gender       0
BMI          23
Children     0
smoking_status 0
location     0
health_insurance_price 2
dtype: int64
```

```
[380]: # creating separate copies for different imputation techniques
df_mean = df.copy()
df_median = df.copy()
df_interpolation = df.copy()
df_ffill = df.copy()
df_bfill = df.copy()
```

```
[381]: # mean imputation
df_mean["age"] = df_mean["age"].fillna(df_mean["age"].mean())
df_mean["BMI"] = df_mean["BMI"].fillna(df_mean["BMI"].mean())
```

```
[382]: df_mean = df_mean.dropna(subset=['health_insurance_price'])
```

```
[383]: # Display missing values
print("Missing values after Mean Imputation:")
print(df_mean.isnull().sum())
```

Missing values after Mean Imputation:

```
age          0
gender       0
BMI          0
Children     0
smoking_status 0
location     0
health_insurance_price 0
dtype: int64
```

```
[384]: # median imputation
df_median["age"] = df_median["age"].fillna(df_median["age"].median())
df_median["BMI"] = df_median["BMI"].fillna(df_median["BMI"].median())
```

```
[385]: df_median = df_median.dropna(subset=['health_insurance_price'])
```

```
[386]: # Display missing values
print("Missing values after Median Imputation:")
print(df_median.isnull().sum())
```

Missing values after Median Imputation:

```
age                0
gender             0
BMI               0
Children          0
smoking_status    0
location          0
health_insurance_price  0
dtype: int64
```

```
[387]: # Display all values are filled
print(df_median.info())
```

```
<class 'pandas.core.frame.DataFrame'>
Index: 1335 entries, 0 to 1337
Data columns (total 7 columns):
#   Column                Non-Null Count  Dtype
---  -
0   age                   1335 non-null  float64
1   gender                1335 non-null  object
2   BMI                  1335 non-null  float64
3   Children              1335 non-null  int64
4   smoking_status        1335 non-null  object
5   location              1335 non-null  object
6   health_insurance_price 1335 non-null  float64
dtypes: float64(3), int64(1), object(3)
memory usage: 83.4+ KB
None
```

i have median to handling missing values in dataset and replace with main dataset (df)

```
[388]: # Interpotaion - Fills missing values using linear interpolation
# median imputation
df_interpolation["age"] = df_interpolation["age"].
    ↪fillna(df_interpolation["age"].interpolate(method="linear"))
df_interpolation["BMI"] = df_interpolation["BMI"].
    ↪fillna(df_interpolation["BMI"].interpolate(method="linear"))
```

```
[389]: df_interpolation = df_interpolation.dropna(subset=['health_insurance_price'])
```

```
[390]: print("Missing values after Interpolation:")
print(df_interpolation.isnull().sum())
```

Missing values after Interpolation:

age	0
gender	0
BMI	1
Children	0
smoking_status	0
location	0
health_insurance_price	0

dtype: int64

```
[391]: # Forward Fill(ffill) - Uses previous values to fill missing values
df_ffill["age"] = df_ffill["age"].ffill()
df_ffill["BMI"] = df_ffill["BMI"].ffill()
```

```
[392]: df_ffill = df_ffill.dropna(subset=['health_insurance_price'])
```

```
[393]: print("Missing values after Forward Fill (ffill):")
print(df_ffill.isnull().sum())
```

Missing values after Forward Fill (ffill):

age	0
gender	0
BMI	1
Children	0
smoking_status	0
location	0
health_insurance_price	0

dtype: int64

```
[394]: # Backward Fill(bfill) - Uses next values to fill missing values
df_bfill["age"] = df_bfill["age"].bfill()
df_bfill["BMI"] = df_bfill["BMI"].bfill()
```

```
[395]: df_bfill = df_bfill.dropna(subset=['health_insurance_price'])
```

```
[396]: print("Missing values after Backward Fill (bfill):")
print(df_bfill.isnull().sum())
```

Missing values after Backward Fill (bfill):

age	0
gender	0
BMI	0
Children	0
smoking_status	0
location	0

```
health_insurance_price    0
dtype: int64
```

I have have selected median to fill null values - df_median

```
[397]: df_median.info()
```

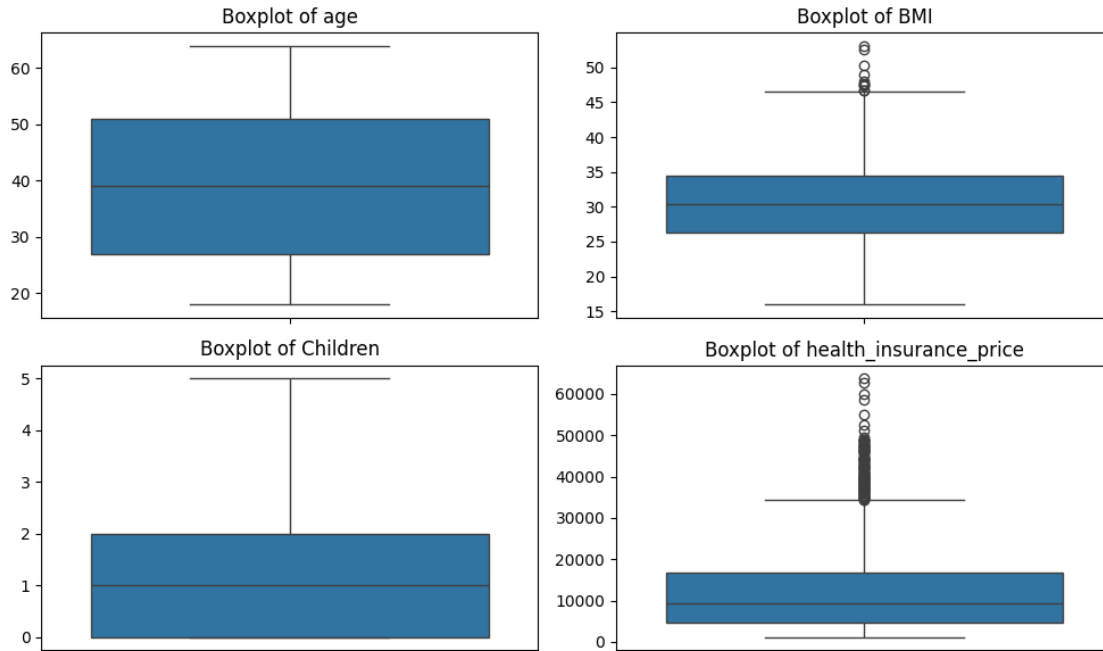
```
<class 'pandas.core.frame.DataFrame'>
Index: 1335 entries, 0 to 1337
Data columns (total 7 columns):
#   Column                Non-Null Count  Dtype
---  -
0   age                   1335 non-null   float64
1   gender                1335 non-null   object
2   BMI                   1335 non-null   float64
3   Children              1335 non-null   int64
4   smoking_status        1335 non-null   object
5   location              1335 non-null   object
6   health_insurance_price 1335 non-null   float64
dtypes: float64(3), int64(1), object(3)
memory usage: 83.4+ KB
```

6 Handling Outliers

```
[398]: from scipy.stats import zscore
```

```
[399]: numerical_features = ['age', 'BMI', 'Children', 'health_insurance_price']

plt.figure(figsize=(10,6))
for i, col in enumerate(numerical_features, 1):
    plt.subplot(2,2,i)
    sns.boxplot(data=df_median, y=col)
    plt.title(f"Boxplot of {col}")
    plt.ylabel("")
plt.tight_layout()
plt.show()
```



```
[400]: # # Detecting Outliers using Z-Score
# z_scores = np.abs(zscore(df.select_dtypes(include=[np.number])))
# outliers_z = (z_scores>3).sum()
# print("Outliers detected using Z-Score:\n",outliers_z)
```

```
[401]: # Select only numeric columns
numeric_df = df_median.select_dtypes(include=["number"])
```

```
[402]: numeric_df
```

```
[402]:
```

	age	BMI	Children	health_insurance_price
0	19.0	30.305	0	16884.92400
1	18.0	33.770	1	1725.55230
2	28.0	33.000	3	4449.46200
3	33.0	22.705	0	21984.47061
4	32.0	28.880	0	3866.85520
...
1333	50.0	30.970	3	10600.54830
1334	18.0	31.920	0	2205.98080
1335	18.0	36.850	0	1629.83350
1336	21.0	25.800	0	2007.94500
1337	61.0	29.070	0	29141.36030

```
[1335 rows x 4 columns]
```

```
[403]: # Compute IQR
Q1 = numeric_df.quantile(0.25)
Q3 = numeric_df.quantile(0.75)
IQR = Q3 - Q1

[404]: # Detect outliers
outliers_iqr = ((numeric_df < (Q1 - 1.5 * IQR))|(numeric_df > (Q3 + 1.5 * IQR))).sum()
print("Outliers detected using IQR:\n",outliers_iqr)
```

```
Outliers detected using IQR:
age          0
BMI          11
Children     0
health_insurance_price  141
dtype: int64
```

6.1 Removing Outliers using IQR (Interquartile Range)

Removes values beyond $1.5 * IQR$ – Good for non-normal data

```
[405]: # Select only numerical columns
numeric_df = df_median.select_dtypes(include=["number"])
```

```
[406]: # Calculate Q1, Q3, and IQR
Q1 = numeric_df.quantile(0.25)
Q3 = numeric_df.quantile(0.75)
IQR = Q3 - Q1
```

```
[407]: # Define lower and upper bounds
lower_bound = Q1 - 1.5 * IQR
upper_bound = Q3 + 1.5 * IQR
lower_bound, upper_bound
```

```
[407]: (age          -9.000000
BMI          14.142500
Children     -3.000000
health_insurance_price  -13066.415245
dtype: float64,
age          87.000000
BMI          46.602500
Children      5.000000
health_insurance_price  34435.221275
dtype: float64)
```

```
[408]: # Identify outliers
outliers_mask = (numeric_df < lower_bound) | (numeric_df > upper_bound)
```

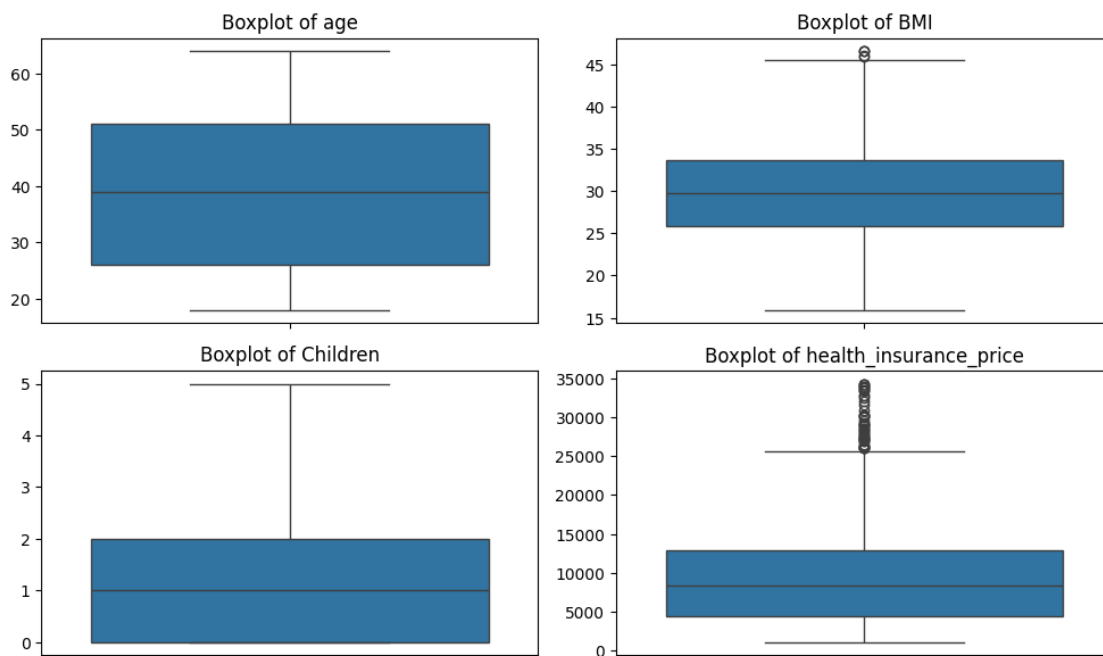
```
[409]: # Filter out outliers
df_cleaned = df_median[~outliers_mask.any(axis=1)]
```

```
[410]: print("Original Data Shape:",df_median.shape)
print("After IQR Outlier Removal:",df_cleaned.shape)
```

Original Data Shape: (1335, 7)
After IQR Outlier Removal: (1186, 7)

```
[411]: # Boxplot for check removing Outliers using IQR (Interquartile Range)
numerical_features = ['age','BMI','Children','health_insurance_price']

plt.figure(figsize=(10,6))
for i, col in enumerate(numerical_features, 1):
    plt.subplot(2,2,i)
    sns.boxplot(data=df_cleaned, y=col)
    plt.title(f"Boxplot of {col}")
    plt.ylabel("")
plt.tight_layout()
plt.show()
```



6.2 Removing Outliers using Z-Score

Removes values with $Z > 3$ – Best for normally distributed data

```
[412]: from scipy.stats import zscore
```



```
[413]: # Compute Z-Scores for numerical columns
z_scores = np.abs(zscore(df_median.select_dtypes(include=[np.number])))
```

```
[414]: print(z_scores)
```

```
[[1.45103968 0.05296611 0.90824519 0.29794375]
 [1.52295867 0.51953905 0.07889681 0.95400578]
 [0.80376878 0.39231568 1.57979994 0.72904939]
 ...
 [1.52295867 1.02843254 0.90824519 0.9619108 ]
 [1.3072017  0.79730544 0.90824519 0.93068414]
 [1.56955786 0.25701918 0.90824519 1.31015191]]
```

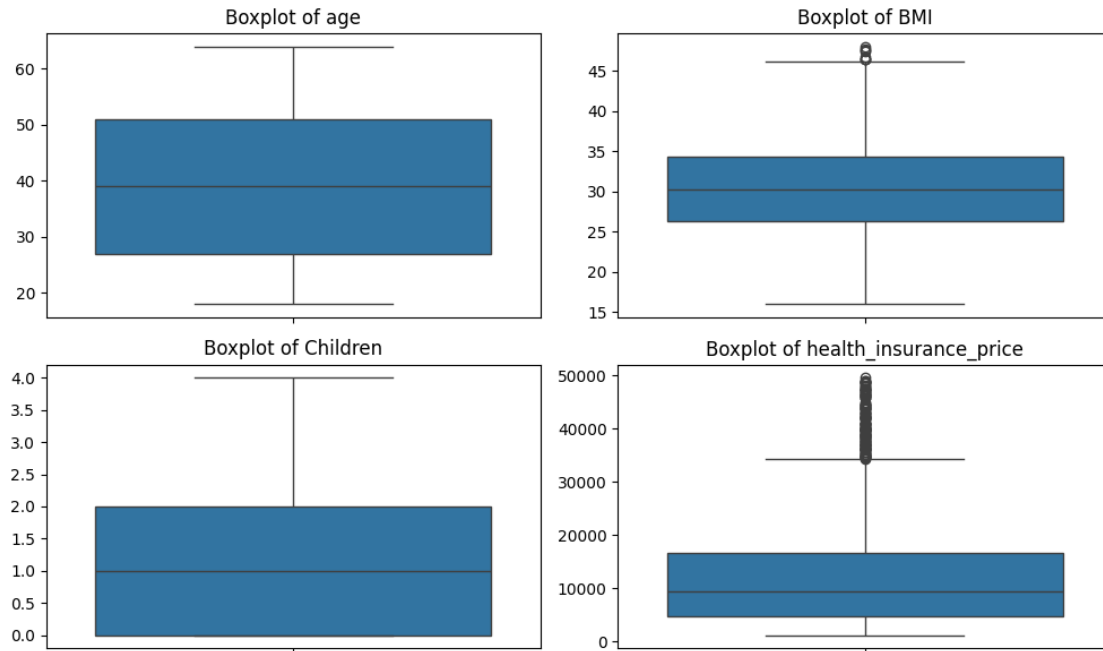
```
[415]: # Keep only data points where Z-score is within ±3
df_z = df_median[(z_scores < 3).all(axis=1)]
```

```
[416]: print("Original Data Shape:",df_median.shape)
print("After Z-Score Outlier Removal:",df_z.shape)
```

```
Original Data Shape: (1335, 7)
After Z-Score Outlier Removal: (1306, 7)
```

```
[417]: # Boxplot check for removed outliers using Z-Score
numerical_features = ['age', 'BMI', 'Children', 'health_insurance_price']

plt.figure(figsize=(10,6))
for i, col in enumerate(numerical_features, 1):
    plt.subplot(2,2,i)
    sns.boxplot(data=df_z, y=col)
    plt.title(f"Boxplot of {col}")
    plt.ylabel("")
plt.tight_layout()
plt.show()
```



6.3 Winsorization (Capping Outliers)

Caps extreme values instead of removing them – Useful when data loss is not acceptable

```
[418]: from scipy.stats.mstats import winsorize
```

```
[419]: # Its store all numerical data
numeric_df
```

```
[419]:
```

	age	BMI	Children	health_insurance_price
0	19.0	30.305	0	16884.92400
1	18.0	33.770	1	1725.55230
2	28.0	33.000	3	4449.46200
3	33.0	22.705	0	21984.47061
4	32.0	28.880	0	3866.85520
...
1333	50.0	30.970	3	10600.54830
1334	18.0	31.920	0	2205.98080
1335	18.0	36.850	0	1629.83350
1336	21.0	25.800	0	2007.94500
1337	61.0	29.070	0	29141.36030

[1335 rows x 4 columns]

```
[420]: numeric_df.isnull().sum()
```

```
[420]: age          0
      BMI          0
      Children     0
      health_insurance_price  0
      dtype: int64
```

```
[421]: df_winsor = df_median.copy()
```

```
[422]: df_winsor
```

```
[422]:      age  gender    BMI  Children  smoking_status  location \
0    19.0  female  30.305         0             yes  southwest
1    18.0   male  33.770         1             no   southeast
2    28.0   male  33.000         3             no   southeast
3    33.0   male  22.705         0             no  northwest
4    32.0   male  28.880         0             no  northwest
...    ...    ...    ...    ...    ...    ...
1333  50.0   male  30.970         3             no  northwest
1334  18.0  female  31.920         0             no  northeast
1335  18.0  female  36.850         0             no  southeast
1336  21.0  female  25.800         0             no  southwest
1337  61.0  female  29.070         0             yes  northwest
```

```
      health_insurance_price
0          16884.92400
1          1725.55230
2          4449.46200
3          21984.47061
4          3866.85520
...          ...
1333          10600.54830
1334          2205.98080
1335          1629.83350
1336          2007.94500
1337          29141.36030
```

```
[1335 rows x 7 columns]
```

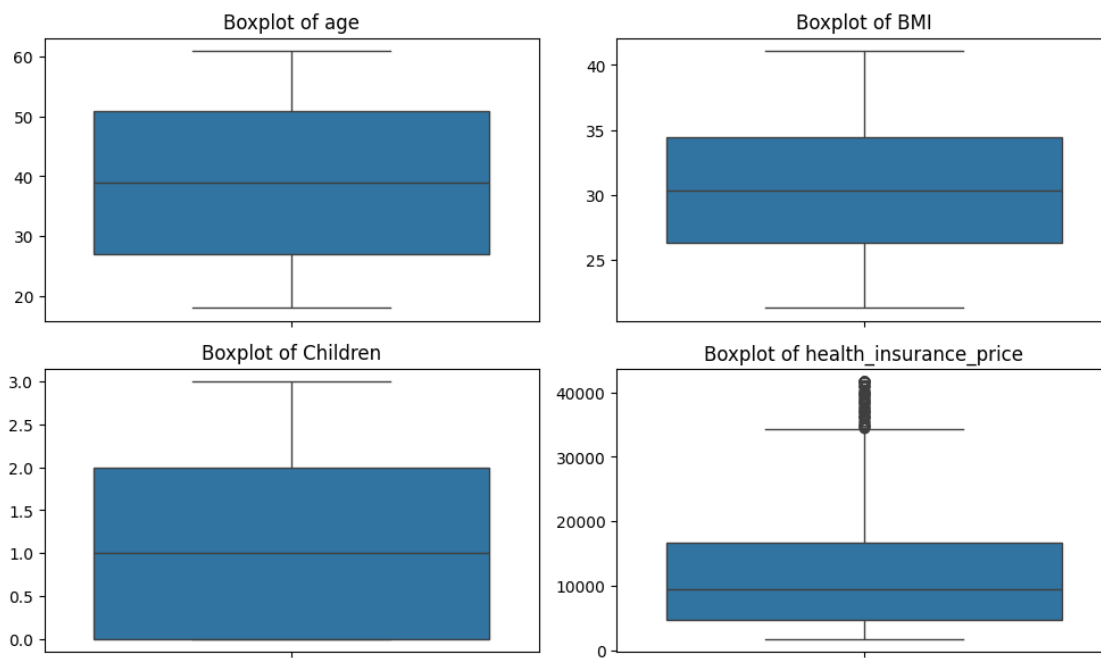
```
[423]: # Apply Winsorization to numerical columns (capping extreme values at 5% and
      ↪ 95%)
      for col in numeric_df.columns:
          df_winsor[col] = winsorize(df_median[col], limits=[0.05,0.05])
```

```
[424]: print("Original Data Shape:",df_median.shape)
      print("After Winsorization (Capping):",df_winsor.shape)
```

Original Data Shape: (1335, 7)
After Winsorization (Capping): (1335, 7)

```
[425]: # using boxplot checking outliers removed method- Winsorization (Capping ↵
        ↵Outliers)
numerical_features = ['age', 'BMI', 'Children', 'health_insurance_price']

plt.figure(figsize=(10,6))
for i, col in enumerate(numerical_features, 1):
    plt.subplot(2,2,i)
    sns.boxplot(data=df_winsor, y=col)
    plt.title(f"Boxplot of {col}")
    plt.ylabel("")
plt.tight_layout()
plt.show()
```



6.4 Scaling

```
[426]: from sklearn.preprocessing import MinMaxScaler, StandardScaler, RobustScaler
```

```
[427]: # Create separate copies of the dataset for each scaling method
df_minmax = df_winsor.copy()
df_standard = df_winsor.copy()
df_robust = df_winsor.copy()
```

```
[428]: # Selecting only numerical columns for scaling
num_cols = df_median.select_dtypes(include=["number"]).columns
```

```
[429]: # Min-Max Scaling (0 to 1 Range)
minmax_scaler = MinMaxScaler()
df_minmax[num_cols] = minmax_scaler.fit_transform(df_minmax[num_cols])
```

```
[430]: df_minmax
```

```
[430]:
```

	age	gender	BMI	Children	smoking_status	location	\
0	0.023256	female	0.455348	0.000000	yes	southwest	
1	0.000000	male	0.630172	0.333333	no	southeast	
2	0.232558	male	0.591322	1.000000	no	southeast	
3	0.348837	male	0.071897	0.000000	no	northwest	
4	0.325581	male	0.383451	0.000000	no	northwest	
...	
1333	0.744186	male	0.488900	1.000000	no	northwest	
1334	0.000000	female	0.536831	0.000000	no	northeast	
1335	0.000000	female	0.785570	0.000000	no	southeast	
1336	0.069767	female	0.228052	0.000000	no	southwest	
1337	1.000000	female	0.393037	0.000000	yes	northwest	

	health_insurance_price
0	0.379066
1	0.000000
2	0.067418
3	0.506867
4	0.052817
...	...
1333	0.221572
1334	0.011193
1335	0.000000
1336	0.006230
1337	0.686227

```
[1335 rows x 7 columns]
```

```
[431]: # Standardize features for model training (Mean=0, Std=1)
standard_scaler = StandardScaler()
df_standard[num_cols] = standard_scaler.fit_transform(df_robust[num_cols])
```

We standardize features using StandardScaler to normalize the input variables before training models.

```
[432]: df_standard
```

```
[432]:
```

	age	gender	BMI	Children	smoking_status	location	\
0	-1.461093	female	-0.048608	-0.955994	yes	southwest	
1	-1.533861	male	0.576850	-0.045036	no	southeast	
2	-0.806177	male	0.437859	1.776879	no	southeast	
3	-0.442334	male	-1.420463	-0.955994	no	northwest	
4	-0.515103	male	-0.305831	-0.955994	no	northwest	
...	
1333	0.794730	male	0.071430	1.776879	no	northwest	
1334	-1.533861	female	0.242911	-0.955994	no	northeast	
1335	-1.533861	female	1.132813	-0.955994	no	southeast	
1336	-1.315556	female	-0.861793	-0.955994	no	southwest	
1337	1.595183	female	-0.271534	-0.955994	yes	northwest	

	health_insurance_price
0	0.337103
1	-0.989647
2	-0.753681
3	0.784413
4	-0.804784
...	...
1333	-0.214135
1334	-0.950469
1335	-0.989647
1336	-0.967840
1337	1.412184

[1335 rows x 7 columns]

```
[433]: # Robust Scaling (Uses Median & IQR - Good for Outliers)
robust_scale = RobustScaler()
df_robust[num_cols] = robust_scale.fit_transform(df_robust[num_cols])
```

```
[434]: print(df_robust)
```

	age	gender	BMI	Children	smoking_status	location	\
0	-0.833333	female	0.000000	-0.5	yes	southwest	
1	-0.875000	male	0.426987	0.0	no	southeast	
2	-0.458333	male	0.332101	1.0	no	southeast	
3	-0.250000	male	-0.936537	-0.5	no	northwest	
4	-0.291667	male	-0.175601	-0.5	no	northwest	
...	
1333	0.458333	male	0.081947	1.0	no	northwest	
1334	-0.875000	female	0.199014	-0.5	no	northeast	
1335	-0.875000	female	0.806531	-0.5	no	southeast	
1336	-0.750000	female	-0.555145	-0.5	no	southwest	
1337	0.916667	female	-0.152187	-0.5	yes	northwest	

```

health_insurance_price
0          0.631453
1         -0.642237
2         -0.415708
3          1.060874
4         -0.464768
...
1333        0.102261
1334       -0.604626
1335       -0.642237
1336       -0.621302
1337        1.663538

```

[1335 rows x 7 columns]

```

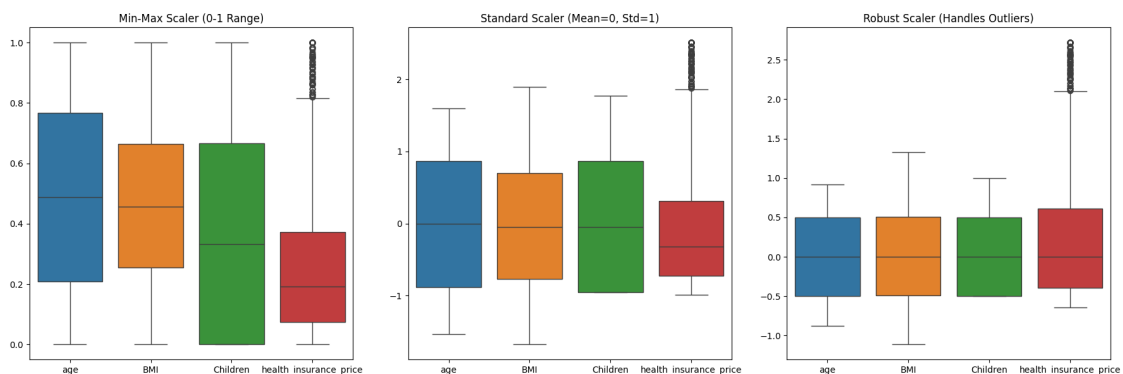
[435]: # Visualizing the differences
fig, axes = plt.subplots(1, 3, figsize=(18, 6))

# Min-Max Scaler
sns.boxplot(data=df_minmax, ax=axes[0])
axes[0].set_title("Min-Max Scaler (0-1 Range)")

# Standard Scaler
sns.boxplot(data=df_standard, ax=axes[1])
axes[1].set_title("Standard Scaler (Mean=0, Std=1)")

# Robust Scaler
sns.boxplot(data=df_robust, ax=axes[2])
axes[2].set_title("Robust Scaler (Handles Outliers)")
plt.tight_layout()
plt.show()

```



6.5 Min-Max Scaler (Left Plot)

1. Scales data to a 0-1 range.
2. The boxplots show values rescaled between 0 and 1.
3. The distribution remains the same, but the scale is compressed.
4. Outliers are still present, which is expected since Min-Max scaling does not handle outliers.

6.6 Standard Scaler (Middle Plot)

1. Standardizes data to have a mean of 0 and a standard deviation of 1.
2. The data is centered around 0.
3. Outliers (values beyond ± 3) are still visible, indicating that extreme values are not removed.
4. This method is sensitive to outliers.

6.7 Robust Scaler (Right Plot)

1. Scales data based on median and IQR, making it resistant to outliers.
2. The boxplots confirm that extreme values are less influential compared to Standard Scaler.
3. This is useful when the dataset contains many outliers.

7 Encoding the Categorical Data

We use Label Encoding to convert categorical variables into numerical format for compatibility with machine learning algorithms.

```
[436]: from sklearn.preprocessing import LabelEncoder
```

```
[437]: # Selecting categorical columns
cat_cols = df_robust.select_dtypes(include=['object']).columns
```

```
[438]: # Encode categorical variables using LabelEncoder
# Initialize the LabelEncoder
la = LabelEncoder()

# Apply label encoding to each categorical column
for col in cat_cols:
    df_robust[col] = la.fit_transform(df_robust[col])
```

```
[439]: # showing first few rows
print("Label Encoded Data:")
print(df_robust.head())
```

Label Encoded Data:

	age	gender	BMI	Children	smoking_status	location	\
0	-0.833333	0	0.000000	-0.5	1	3	
1	-0.875000	1	0.426987	0.0	0	2	
2	-0.458333	1	0.332101	1.0	0	2	
3	-0.250000	1	-0.936537	-0.5	0	1	
4	-0.291667	1	-0.175601	-0.5	0	1	


```

    health_insurance_price
0          0.631453
1         -0.642237
2         -0.415708
3          1.060874
4         -0.464768

```

```
[440]: print("\nShape of the DataFrame (rows, columns):")
print(df_robust.shape)
```

```

Shape of the DataFrame (rows, columns):
(1335, 7)

```

8 Data Splitting (Train-Test Split)

```
[441]: from sklearn.model_selection import train_test_split
```

```
[442]: # Define features (X) and target variable (y)
X = df_robust.drop(columns=['health_insurance_price']) # features
y = df_robust['health_insurance_price'] # target variable
```

```
[443]: print(X)
```

```

      age  gender      BMI  Children  smoking_status  location
0  -0.833333      0  0.000000      -0.5             1          3
1  -0.875000      1  0.426987       0.0             0          2
2  -0.458333      1  0.332101       1.0             0          2
3  -0.250000      1 -0.936537      -0.5             0          1
4  -0.291667      1 -0.175601      -0.5             0          1
...     ...    ...      ...      ...             ...      ...
1333  0.458333      1  0.081947       1.0             0          1
1334 -0.875000      0  0.199014      -0.5             0          0
1335 -0.875000      0  0.806531      -0.5             0          2
1336 -0.750000      0 -0.555145      -0.5             0          3
1337  0.916667      0 -0.152187      -0.5             1          1

```

```
[1335 rows x 6 columns]
```

```
[444]: print(y)
```

```

0          0.631453
1         -0.642237
2         -0.415708
3          1.060874
4         -0.464768

```

```
...
1333    0.102261
1334   -0.604626
1335   -0.642237
1336   -0.621302
1337    1.663538
Name: health_insurance_price, Length: 1335, dtype: float64
```

```
[445]: # Perform train-test split (Stratified for imbalanced data)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
↳ random_state=42)
```

We split the balanced dataset into training and testing sets for model evaluation.

```
[446]: print(X.shape, X_train.shape, X_test.shape)
```

```
(1335, 6) (1068, 6) (267, 6)
```

9 Linear Regression Model Training

```
[447]: from sklearn.linear_model import LinearRegression
```

```
[448]: # loading linear regression model
lr_model = LinearRegression()
```

```
[449]: # Train the model
lr_model.fit(X_train, y_train)
```

```
[449]: LinearRegression()
```

```
[450]: # Make predictions on test data
y_pred_lr = lr_model.predict(X_test)
```

```
[451]: print("Linear Regression Performance")
print(f"Accuracy Score:", lr_model.score(X_test, y_test)*100)
```

```
Linear Regression Performance
Accuracy Score: 82.76836754373782
```

- The Linear Regression model achieved an accuracy score of 82.77%, indicating that it correctly predicted the outcomes approximately 83% of the time.
- The model shows good generalization, but further improvements can be made through non-linear relationships, or by exploring other regression models

10 Model Evaluation for Linear Regression

```
[452]: from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score
```

```
[453]: # Model Evaluation:
print("MAE:", mean_absolute_error(y_test, y_pred_lr))
print("MSE:", mean_squared_error(y_test, y_pred_lr))
print("RMSE:", mean_squared_error(y_test, y_pred_lr))
print("R2 Score:", r2_score(y_test, y_pred_lr))
```

MAE: 0.32899403640691405

MSE: 0.1961106631835952

RMSE: 0.1961106631835952

R² Score: 0.8276836754373782

- The model fits the data well, especially for $R^2 > 0.8$ is strong.
- The errors (MAE and RMSE) are relatively low, suggesting the predictions are fairly accurate on average.
- There may still be some outliers or non-linear patterns not captured (hinted by MSE being ~0.196).

```
[454]: for name, coef in zip(X.columns, lr_model.coef_):
        print(f"{name}: {coef}")
```

age: 0.49268972582183324

gender: -0.022411843103699217

BMI: 0.20642310929697338

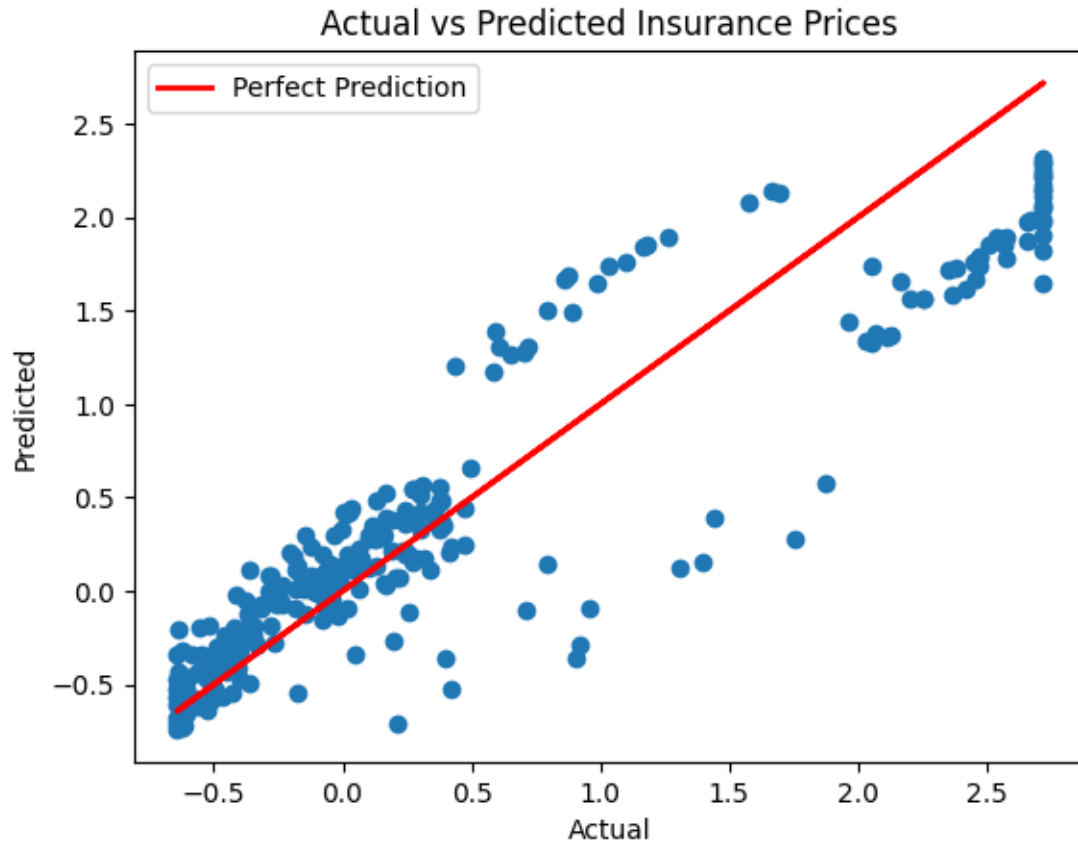
Children: 0.0930957167446882

smoking_status: 1.8326669368332433

location: -0.02249844844852761

- smoking_status is by far the most influential predictor, which makes sense in real-world insurance pricing.
- age and BMI also have clear, positive effects.
- gender and location appear statistically insignificant in this model.

```
[455]: plt.scatter(y_test, y_pred_lr)
plt.plot(y_test, y_test, color='red', linewidth=2, label='Perfect Prediction')
plt.xlabel("Actual")
plt.ylabel("Predicted")
plt.title("Actual vs Predicted Insurance Prices")
plt.legend()      # This shows the label for the red line
plt.show()
```



- The red diagonal line represents the perfect prediction line (i.e., where Predicted = Actual).
- The blue scatter points represent the predicted values from your model against the actual test values.
- Many points are close to the red line, indicating the model gets many predictions reasonably accurate.
- These might be outliers or cases where key features didn't generalize well.

11 Random Forest Regressor Model Training

```
[456]: from sklearn.ensemble import RandomForestRegressor  
from sklearn.datasets import make_regression
```

```
[457]: # Loading Random Forest Regressor Model  
rf_model = RandomForestRegressor(max_depth=2, random_state=0)
```

```
[458]: # Train the model  
rf_model.fit(X_train, y_train)
```

```
[458]: RandomForestRegressor(max_depth=2, random_state=0)
```

```
[459]: y_pred_rf = rf_model.predict(X_test)
```

```
[460]: print("Random Forest Regression Performance")
print(f"Accuracy Score:", rf_model.score(X_test, y_test)*100)
```

Random Forest Regression Performance
Accuracy Score: 89.09440164404525

- The Random Forest Regressor model explains approximately 89% of the variance in the target variable on the test data. This indicates a strong performance, suggesting that the model is able to accurately predict the target variable for most data points.

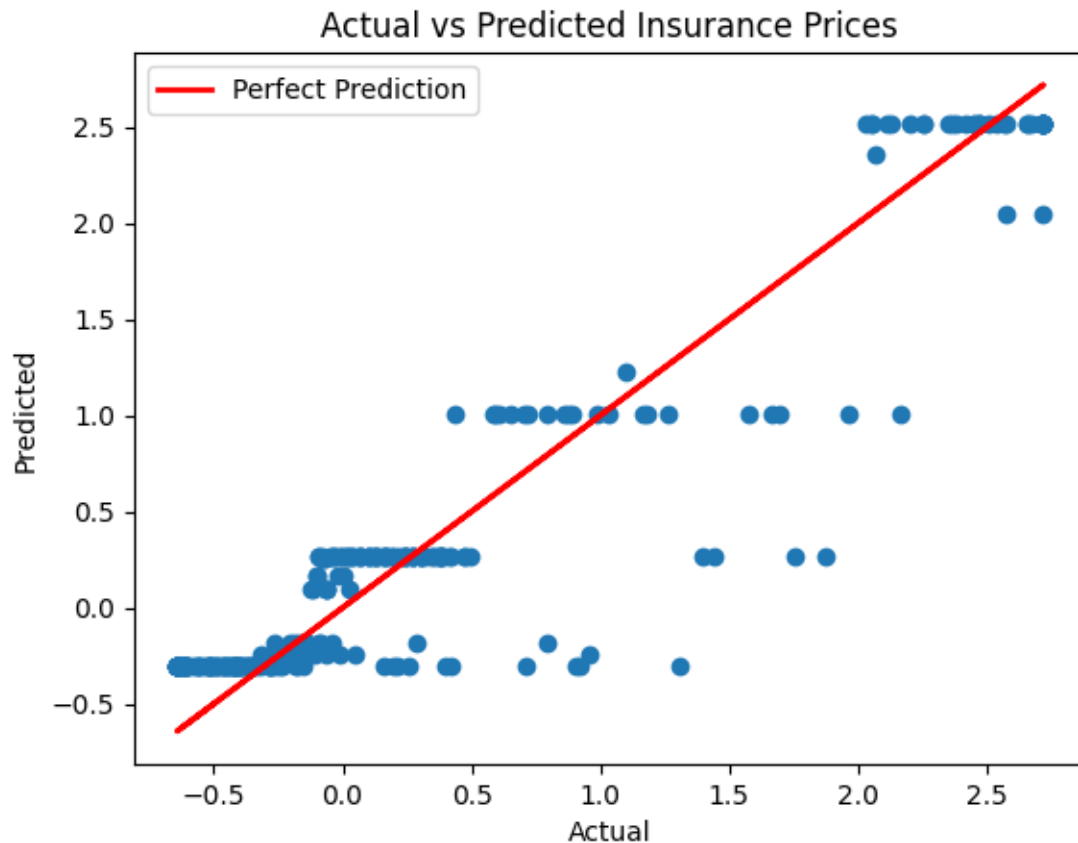
12 Model Evaluation for Random Forest Regressor

```
[461]: # Model Evaluation:
print("MAE:", mean_absolute_error(y_test, y_pred_rf))
print("MSE:", mean_squared_error(y_test, y_pred_rf))
print("RMSE:", mean_squared_error(y_test, y_pred_rf))
print("R2 Score:", r2_score(y_test, y_pred_rf))
```

MAE: 0.23782627921808241
MSE: 0.12411500369618084
RMSE: 0.12411500369618084
R² Score: 0.8909440164404525

- Your Random Forest model performs well with low error values and a high R² score.
- The predictions are generally close to the actual values.

```
[462]: plt.scatter(y_test, y_pred_rf)
plt.plot(y_test, y_test, color='red', linewidth=2, label='Perfect Prediction')
plt.xlabel("Actual")
plt.ylabel("Predicted")
plt.title("Actual vs Predicted Insurance Prices")
plt.legend() # This shows the label for the red line
plt.show()
```



- The Random Forest model predicts insurance prices effectively and consistently, as seen by the close clustering around the perfect prediction line.

13 Decision Tree Regressor Model Training

```
[463]: from sklearn.tree import DecisionTreeRegressor
```

```
[464]: # loading and training decision tree regression model
dt_model = DecisionTreeRegressor(random_state=0)
dt_model.fit(X_train, y_train)
```

```
[464]: DecisionTreeRegressor(random_state=0)
```

```
[465]: # Make predictions on test data
y_pred_dt = dt_model.predict(X_test)
```

```
[466]: print("Decision Tree Regression Performance")
print(f"Accuracy Score:", dt_model.score(X_test, y_test)*100)
```

Decision Tree Regression Performance
Accuracy Score: 78.26582370054028

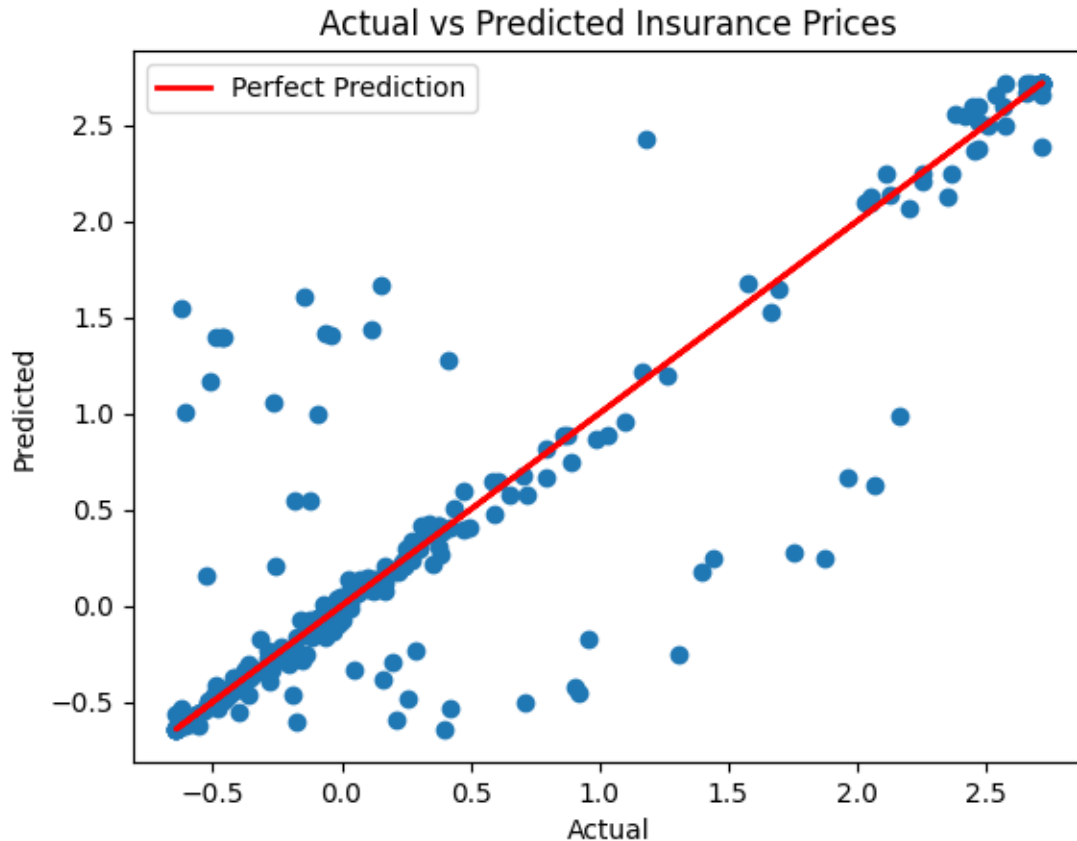
- The Decision Tree Regression model performs reasonably well, achieving an accuracy score of 78.27%.
- However, it is outperformed by the Random Forest model, which offers better accuracy and generalization.

14 Model Evaluation for Decision Tree Regressor

```
[467]: # Model Evaluation:
print("MAE:", mean_absolute_error(y_test, y_pred_dt))
print("MSE:", mean_squared_error(y_test, y_pred_dt))
print("RMSE:", mean_squared_error(y_test, y_pred_dt))
print("R2 Score:", r2_score(y_test, y_pred_dt))
```

MAE: 0.21669262956747398
MSE: 0.24735344945726542
RMSE: 0.24735344945726542
R² Score: 0.7826582370054028

```
[468]: plt.scatter(y_test, y_pred_dt)
plt.plot(y_test, y_test, color='red', linewidth=2, label='Perfect Prediction')
plt.xlabel("Actual")
plt.ylabel("Predicted")
plt.title("Actual vs Predicted Insurance Prices")
plt.legend() # This shows the label for the red line
plt.show()
```



- The scatter plot suggests that the Decision Tree Regressor has moderate predictive performance.
- It works well in some areas (especially low-to-mid values) but lacks generalization for higher or more complex data points.
- Overfitting is likely in some regions, which is a known weakness of decision trees when not pruned or tuned properly.

15 Model Comparison

summarize performance metrics in a comparison table and visualize them using grouped bar charts to help identify the best-performing model.

```
[469]: # Import performance metrics
from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score
import numpy as np

# Make predictions using all three trained models
y_pred_lr = lr_model.predict(X_test)
y_pred_rf = rf_model.predict(X_test)
```



```

y_pred_dt = dt_model.predict(X_test)

# Calculate metrics, Store model performance metrics in a dictionary
metrics = {
    'Model': ['Linear Regression', 'Random Forest Regressor', 'Decision Tree_
↪Regressor'],
    'MAE': [
        mean_absolute_error(y_test, y_pred_lr),
        mean_absolute_error(y_test, y_pred_rf),
        mean_absolute_error(y_test, y_pred_dt)
    ],
    'MSE': [
        mean_squared_error(y_test, y_pred_lr),
        mean_squared_error(y_test, y_pred_rf),
        mean_squared_error(y_test, y_pred_dt)
    ],
    'RMSE': [
        np.sqrt(mean_squared_error(y_test, y_pred_lr)),
        np.sqrt(mean_squared_error(y_test, y_pred_rf)),
        np.sqrt(mean_squared_error(y_test, y_pred_dt))
    ],
    'R² Score': [
        r2_score(y_test, y_pred_lr),
        r2_score(y_test, y_pred_rf),
        r2_score(y_test, y_pred_dt)
    ]
}

# Display results
results_df = pd.DataFrame(metrics)
print("Model Performance Comparison:\n")
print(results_df)

# Identify and print the model with the best accuracy
best_model = results_df.loc[results_df['R² Score'].idxmax()]
print(f"\nBest Performance Model: {best_model['Model']} with R² Score:
↪{best_model['R² Score']:.4f}")

```

Model Performance Comparison:

	Model	MAE	MSE	RMSE	R² Score
0	Linear Regression	0.328994	0.196111	0.442844	0.827684
1	Random Forest Regressor	0.237826	0.124115	0.352300	0.890944
2	Decision Tree Regressor	0.216693	0.247353	0.497346	0.782658

Best Performance Model: Random Forest Regressor with R² Score: 0.8909

16 Visual Comparison of Model Performance

```
[470]: # Set up bar chart positions and width
bar_width = 0.2
models = results_df['Model']
x = np.arange(len(models))

# Create a new figure for the bar chart
fig, ax = plt.subplots(figsize=(10, 5))

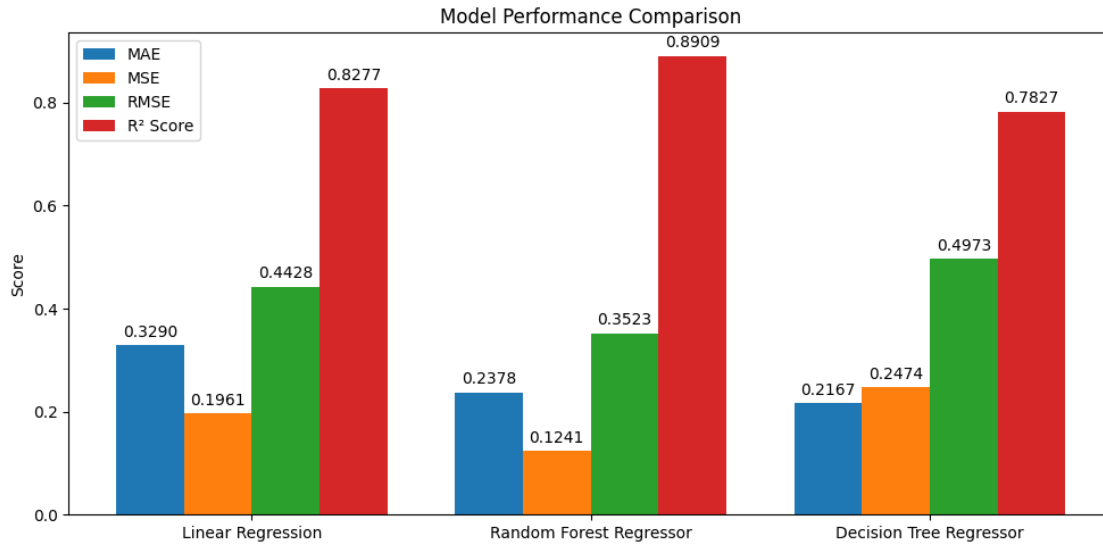
# Plot bars for each metric
bars1 = ax.bar(x - 1.5*bar_width, results_df['MAE'], width=bar_width,
               label='MAE')
bars2 = ax.bar(x - 0.5*bar_width, results_df['MSE'], width=bar_width,
               label='MSE')
bars3 = ax.bar(x + 0.5*bar_width, results_df['RMSE'], width=bar_width,
               label='RMSE')
bars4 = ax.bar(x + 1.5*bar_width, results_df['R2 Score'], width=bar_width,
               label='R2 Score')

# Function to add value labels on top of bars
def add_labels(bars):
    for bar in bars:
        height = bar.get_height()
        ax.annotate(f'{height:.4f}',
                    xy=(bar.get_x() + bar.get_width() / 2, height),
                    xytext=(0, 3), # vertical offset
                    textcoords="offset points",
                    ha='center', va='bottom')

# Add labels to all bars
for bars in [bars1, bars2, bars3, bars4]:
    add_labels(bars)

# Set x-axis labels and chart title
ax.set_xticks(x)
ax.set_xticklabels(models)
ax.set_ylabel('Score')
ax.set_title('Model Performance Comparison')
ax.legend()

# Adjust layout and display the plot
plt.tight_layout()
plt.show()
```



17 Results:

- Random Forest delivered the best performance, striking a strong balance between MSE and RMSE.
- Reduced false negatives, which are critical in fraud detection.
- Random Forest Regressor outperformed the other models, achieving the highest R² score of 0.8909 and the lowest error rates (MAE: 0.2378, MSE: 0.1241, RMSE: 0.3523), indicating strong predictive accuracy.
- Linear Regression showed a decent performance with an R² score of 0.8277 but higher error rates compared to Random Forest.
- Decision Tree Regressor had the lowest MAE (0.2167) but a lower R² score (0.7827) and the highest RMSE (0.4973), suggesting less consistency in predictions.

18 Conclusion:

- Random Forest Regressor is the most suitable model for this problem, offering the best balance between accuracy and error minimization.
- The model revealed that smoking status, age, and BMI are the most influential factors in determining insurance costs.
- Ensemble techniques like Random Forest help reduce variance and improve generalization, making them more robust than individual decision trees or linear models.